

Project 4

Steg-o-matic

Time due: 9 PM Thursday, March 12th

Introduction.....	2
Anatomy of a Steganography System.....	2
Encoding a Message	2
Decoding a Message	3
What Do You Need to Do?	4
What Will We Provide?	6
The HTTP Class (aka But I don't know how to use C++ to access the Internet!)	6
Details: The Classes You Must Write.....	7
HashTable Class Template	8
The HashTable Class Template: Details	13
BinaryConverter Class	17
The BinaryConverter Class: Details	19
Compressor Class.....	20
The Compressor Class: Details	26
Steg Class.....	27
The Steg Class: Details	30
WebSteg Class	30
The WebSteg Class: Details.....	31
Requirements and Other Thoughts	31
What to Turn In.....	32
Grading	33

Before writing a single line of code, you must first read AND THEN RE-READ the Requirements and Other Thoughts section.

Introduction

The NachenSmall Software Corporation has been contacted by the Paranoid Students Users Group (PSUG), whose members believe that they're constantly being watched by "the man," to create a tool that helps them to communicate with each other in a secret manner. More specifically, the PSUG would like NachenSmall's team to build a steganography system capable of embedding hidden secret messages within ordinary web pages.

Steganography is the practice of embedding a digital message or file within another "carrier" digital medium. For example, you could encode a secret message like "Meet me in Ackerman at noon" by making certain very minor color changes in the pixels of a JPG image. The idea behind such a steganographic encoding is that someone viewing the carrier medium (e.g., the JPG) would not even know that it contains a secret message, so the carrier medium can be freely distributed or posted on the Internet for all to see. Yet the recipient, possessing the knowledge that the carrier medium actually contains an embedded message, and equipped with the proper tools, can extract that message and read it. Steganography has been used for years by spies, terrorist organizations, and almost certainly by paranoid individuals like those in PSUG¹.

So, in your last project for CS32 Winter 2015, your goal is to build a simple set of C++ classes that can be used to implement a steganography system. If you're able to prove to NachenSmall's reclusive and bizarre CEO, Carey Nachenberg, that you have the programming skills to build the simple steganography tool described in this specification, he'll hire you to build the complete project, and you'll be famous... at least in the paranoid student community.

Anatomy of a Steganography System

Steganography tools have two modes of operation:

1. Encoding a message/file within a carrier medium like a web page or JPG, and
2. Decoding or extracting a previously encoded message/file from a carrier medium

Encoding a Message

The steganographic tool that you will build encodes a secret message within a carrier web page, using the following steps:

¹ In fact, steganography is a key element in Carey Nachenberg's upcoming techno-thriller novel, *The Florentine Deception* – get your copy on April 14, and Carey will sign it!

1. It compresses the secret message, like “Meet me in Ackerman at noon”, into a more compact numeric form, like 7234 25115 3451 58345 45636. We will describe exactly how to do this in the specification below.
2. It converts the resulting series of numbers into a binary form, a string of 1 and 0 characters that represents the compressed message.
3. It converts this long string of 1s and 0s into a string made up of tabs and spaces (which in this specification we will make visible by representing them with - and _ respectively), so “101101” might be converted into “- _ _ -”.
4. It retrieves a web page from the Internet that the user has chosen to use as a carrier file to hold the secret message.
5. It splits the web page into separate lines (where each line is terminated by a newline character).
6. It strips all existing whitespace (tabs and spaces) from the end of each line of the original web page, so, for example the line “<h1>Welcome to my web page</h1> - _ -” (the line ends with space tab space tab) would become “<h1>Welcome to my web page</h1>”.
7. It splits up the tabs-and-spaces string that encodes the secret message into substrings, one for each line of the web page, and appends each substring onto the end of a stripped line of that web page.
8. It joins the all of the updated lines of the web page into one string, which is now the carrier of the encoded message. This string might be saved to a file to be served up by a web server.

Someone looking at the newly-created web page in a browser will see nothing suspicious. The rare person who might look at the HTML source of the web page might not notice the extra tabs and spaces at the ends of lines.

Decoding a Message

To extract and decode a secret message that was previously hidden within a web page, the steganographic tool will use the following steps:

1. It retrieves the carrier web page from the Internet. This carrier HTML file was presumably posted on the Internet by the user that created it, and it holds a previously-encoded message.
2. It breaks up the file into different lines (where each line is terminated by a newline character).
3. From each line, it extracts all of the whitespace (tabs and spaces) from the end of the line – this is where the secret message was added during the encoding step.
4. It concatenates all of the whitespace extracted from all of the lines together into a single long string of tabs and spaces.
5. It converts this tabs-and-spaces string into a string of 1 and 0 characters.
6. It converts the resulting string in binary form into a series of numbers, like 7234 25115 3451 58345 45636.

7. It uses a decompression algorithm to expand these numbers back into the original secret message, e.g., “Meet me in Ackerman at noon”.
8. It returns the original secret message to the user.

What Do You Need to Do?

So, at a high level, what do you need to build?

You’ll be building five complete classes and two new stand-alone functions, described at a high level below. Detailed specs follow in the later sections.

You need to build a hash table class template *HashTable*:

1. You need to be able to create a new *open hash table* with a specified number of buckets and a specified capacity.
2. You need to be able to add new associations (each association is a key→value mapping) to the hash table, e.g., 1234→“meet”, or “Ackerman”→45235.
3. You need to be able to efficiently search for an association in the hash table given a key.
4. You need to be able to determine if the hash table has reached its maximum capacity.
5. You need to be able to efficiently mark an association in the hash table as having been recently written (i.e., added or modified).
6. You need to be able to efficiently discard the least recently written association from the hash table. (This will be used if the hash table is full and its client needs to remove an association from the table to make room for a new one.

You need to build a class *BinaryConverter*:

1. Given a vector of unsigned short integers, this class must be able to convert the vector into a string containing just tabs and spaces, representing the binary form (1s and 0s) of the integers in the vector.
2. Given a string containing just tabs and spaces, this class must be able convert this back into its binary equivalent, and from there, back into a vector of unsigned short integers.

Since you haven’t learned binary encoding yet, we will also provide you with a few functions to help you convert a single unsigned short into a binary string, and vice versa.

You need to build a class *Compressor*:

1. This class enables you to compress a secret message, like “Meet me in Ackerman at noon”, into a series of numbers, like 7234 25115 3451 58345 45636.

2. This class enables you to decompress a previously-created series of numbers, like 7234 25115 3451 58345 45636, back into the original secret message, “Meet me in Ackerman at noon”.

Your *Compressor* class **must** produce the same result as our prescribed compression and decompression algorithms exactly, or your program will not function properly with our test tools and you will receive zero credit on this part of the project. Your *Compressor* class must use your *HashTable* class template when compressing and decompressing – we’ll tell you exactly how that works in the sections below.

You need to build a class *Steg*:

1. This class allows its client to hide a secret message, like “Meet me in Ackerman at noon”, within a string. This involves:
 - a. Compressing the secret message using the *Compressor* class into a sequence of unsigned short integers .
 - b. Converting those integers into their binary equivalent, and from that into a corresponding string of tabs and spaces.
 - c. Stripping existing tabs and spaces from the end of each line of the string.
 - d. Appending the tabs and spaces that represent the secret message onto the ends of the lines of the stripped lines.
 - e. Combining the lines and returning the string that contains the hidden message.
2. This class allows its client to extract and decode a secret message that was previously encoded inside a carrier string. This involves:
 - a. Extracting the tabs and spaces from the end of each line of a string that previously had a message hidden in it, and creating one long string of tabs and spaces.
 - b. Converting the tabs-and-spaces string to its binary equivalent, and then converting the binary string back to the original sequence of integers
 - c. Using the *Compressor* class to decompress these integers into the original secret message.
 - d. Returning the revealed secret message string.

You need to build a class *WebSteg*:

1. This class allows its client to specify a URL for a web page and a secret message, and produces a modified version of that page that contains the hidden message.
 - a. The class must connect over the Internet to retrieve the web page specified by the URL (using the HTTP interface that we provide).
 - b. The class must use the *Steg* class to hide the provided secret message within the page.
 - c. The class must return a string with the content of the resulting page.
2. This allows its client to specify a URL for a web page that already contains a secret message embedded within it, returns the decoded secret message.

- a. The class must connect over the Internet to retrieve the web page specified by the URL. (This page contains the steganographically-encoded secret message.)
- b. The class must use the *Steg* class to extract and decode the secret message that is embedded within the page
- c. The class must return a string with the decoded secret message.

You need to create two overloaded hash functions, each named `computeHash()`:

1. One *computeHash()* function must take a string argument and compute an unsigned integer hash value from that string.
2. The other *computeHash()* function must take an unsigned short integer and compute an unsigned integer hash value for this unsigned short integer. (This should be a pretty simple function!)

What Will We Provide?

We'll provide a header file, *provided.h*, that contains the interfaces for the *BinaryConverter*, *Compressor*, *Steg*, and *WebSteg* classes.

We'll provide some helper code for you in *BinaryConverter.cpp* that converts unsigned short values to binary, and vice versa.

We'll provide an *HTTP* class that can be used to download a web page from a web server on the Internet (e.g., from *http://reddit.com*). If you specify the URL for a page, it will download the contents of the page and place them into a string.

We'll provide a simple *main.cpp* file that lets you test your overall steganography implementation.

The HTTP Class (aka But I don't know how to use C++ to access the Internet!)

Oh, we knew you were going to say that! Such a whiner! But wouldn't you like to learn how to write a program that interacts with other computers over the Internet? We thought so. So we're going to provide you with a reasonably functional Internet HTTP interface that is capable of downloading pages off of the Internet for you. HTTP is the protocol used by web browsers to download web pages from servers on the Internet into your browser.

When you use our interface, you don't have to worry about the details of how to communicate over the Internet yourself. Of course, if you want to see how our interface works, you're welcome to do so... and before you know it, you'll be forming your own

start-up Internet company to compete against Google². Our HTTP interface's primary public function (`get`) is as easy to use as this:

```
#include "http.h"

int main()
{
    string url = "http://en.wikipedia.org/wiki/Bald";
    string page; // to hold the contents of the web page

    // The next line downloads a web page for you. So easy!
    if (HTTP().get(url, page))
        cout << page; // prints the page's data out
    else
        cout << "Error fetching content from URL " << url << endl;
    ...
}
```

Note that you don't need to declare an HTTP variable. The call above looks as if it calls a function named `HTTP`, then calls a `get()` member function on what it returns.

A challenge when testing a program that analyzes the contents of web pages is that you have no control over those contents. Our HTTP interface lets you set up a pseudo-web of pages with URLs and contents of your choosing:

```
int main()
{
    HTTP().set("http://a.com", "This is a test page.");
    HTTP().set("http://b.com", "Here is another.");
    HTTP().set("http://c.com", "<html>Everyone loves CS 32</html>");
    string page;
    if (HTTP().get("http://b.com", page))
        cout << page << endl; // writes Here is another.
}
```

You call `set()` to associate a URL with a string. From that point on, calling `get()` with that URL will retrieve that string. (Once you call `set()`, `get()` will no longer retrieve pages from the real web; it will instead consult only the pages that you installed with `set()`.)

Details: The Classes You Must Write

You must write correct versions of the following classes to obtain full credit on this project. Your classes must work correctly with our provided code, and you must **not** modify our provided code to make it work with your code. Doing so will result in a **zero score** on that part of the project.

² By your agreeing to use our HTTP code for Project 4, this license entitles NachenSmall to a 20% cut of all profits.

HashTable Class Template

You must write a class template named *HashTable* that lets a client associate items of a key type with items of a (usually different) value type, with the ability to look up items by key efficiently. Your implementation **must** use an *open hash table*.

A HashTable whose key type is string and value type is int, for example, lets a client add A number of key→value associations (e.g., “carey”→1, “david”→2, “lily”→3, etc.), and allows the client to search using a particular key and find the associated value.

When you construct a new *HashTable* object, you must specify both the number of *buckets* in the hash table, as well as its *capacity*, the maximum number of associations that can be stored in the table. If a *HashTable* object is filled to its capacity and the client tries to add an association with a key that’s not already present in the table, then the addition will fail. However, even if the hash table is at capacity, the client may still update an association that has an existing key in the table by replacing its associated value with another value.

Of course, when you insert an association into a hash table, you need to be able to compute the hash of the key for that association. Therefore, in addition to defining your *HashTable* class template, you must also define a non-member function named *computeHash* for each type of key you wish to put into your hash table. Look for examples of these functions in the code snippets below.

For example, someone could use your *HashTable* class template to associate the following key strings with integer values:

```
#include "HashTable.h"

// here's a computeHash() function that hashes a C++ string to an
// unsigned int hash value

unsigned int computeHash(string s)
{
    unsigned int hashValue;

    // Here you'll write code that computes an unsigned int hash
    // value from a C++ string, and returns it. A HashTable whose
    // key type is string will call this function as part of
    // determining a bucket number.

    return hashValue;
}

void argh()
{
    // create a hash table object that has 100 buckets
    // and has a capacity of 200 associations

    HashTable<string, int> nameToAge(100, 200);
```



```

    if ( ! nameToAge.set("Carey", 43) )
        cout << "Error associating Carey with 43!\n";
    else
        cout << "Carey is mapped to 43\n";

    if ( ! nameToAge.set("David", 97) )
        cout << "Error associating David with 97!\n";
    else
        cout << "David is mapped to 97\n";

    int age;
    if ( nameToAge.get("Carey", age) )
        cout << "Carey is associated with age: " << age;
    else
        cout << "Unable to find an association for Carey\n";

    ...
}

```

Here's another example:

```

// hashes a Rectangle value to an unsigned int hash value
unsigned int computeHash(Rectangle r)
{
    // Here's an example of how we might compute a hash for a
    // rectangle.

    return r.getX() + r.getY() + r.getWidth() + r.getHeight();
}

void hashABunchOfRects()
{
    Rectangle a(2, 5, 10, 20);
    Rectangle b(0, 0, 30, 40);
    Rectangle c(10, 7, 100, 15);

    HashTable<Rectangle,unsigned int> rectToArea;

    rectToArea.set(a, a.getArea()); // maps a to its area
    rectToArea.set(b, b.getArea()); // maps b to its area
    rectToArea.set(c, c.getArea()); // maps c to its area

    unsigned int area;
    if ( rectToArea.get(b, area) )
        cout << "The area of rectangle b is: " << area;
}

```

Note that the *HashTable* class must maintain no more than one association with a particular key, so if a client associates “Joe”→5 and then asks to associate “Joe”→17, there will be only one association in the table with the key “Joe”, and that association will have the value 17. The earlier mapping to 5 was replaced by the later mapping to 17.

Unfortunately for you, you're not just going to implement a conventional open hash table for this problem. There's a special requirement that will make implementing this class a bit (a lot?) more challenging.

Your version of *HashTable* must be able to track how each association in the hash table was written (i.e., added or modified). So for instance, imagine that we add a bunch of associations into a *HashTable* instance over time:

```
void argh()
{
    HashTable<string, int> nameToAge(100, 200);

    // add a bunch of associations in this order
    nameToAge.set("Carey", 43);
    nameToAge.set("David", 97);
    nameToAge.set("Lily", 18);
    nameToAge.set("Sally", 22);
    nameToAge.set("David", 55);
}
```

The resulting *HashTable* object will hold:

Carey → 43
Lily → 18
Sally → 22
David → 55

David's association will be the most recently written association in the *HashTable*, with Sally's being the next most recently written one, then Lily's, and then Carey's. Your *HashTable* class template must track the order of how recently each association has been written within the table.

Specifically, any time you *add* or *update* an association, that association becomes the most recently written one, and all earlier-written items fall down one position in the ranking.

There is one exception to this rule. When adding an association with a key not already in the hash table using the *set()* method, you may designate the item as being "permanent." Permanent items ***must not*** be tracked in your hash table's recently written association list. The client of your class may designate an item as permanent by passing a third argument value of *true* to your *set()* method when adding a new association. (A third argument of *false*, just like no third argument at all, means the association is not permanent, and its recency of being written must be tracked.)

```
void argh()
{
    HashTable<string, int> nameToAge(100, 200);

    // add a bunch of associations in this order
    nameToAge.set("Carey", 43, true);    // permanent
}
```

```

    nameToAge.set("Lily", 18);
    nameToAge.set("Sally", 22);
    nameToAge.set("David", 97, true);    // permanent
    nameToAge.set("Diego", 17);
    nameToAge.set("Sally", 23);
}

```

In the above example, both Carey and David’s associations would be designated as permanent in your hash table and must not be tracked in your recently-written list. Lily’s, Sally’s and Diego’s associations would be tracked in your list, with Sally’s being the most recently written, then Diego’s, and finally Lily’s.

You must also enable the client of your *HashTable* class to be able to manually “touch” an existing non-permanent item that is already in the *HashTable* without actually changing its value, to indicate that it should be brought to the top of the most recently written list:

```

void argh()
{
    HashTable<string, int> nameToAge(100,200);

    nameToAge.set("Cameron", 25);
    nameToAge.set("Lily", 18);
    nameToAge.set("Timothy", 43, true);
    nameToAge.set("Sally", 22);
    nameToAge.set("Mikey", 28);

    // This will move Cameron to the top of the most-recently
    // written list
    nameToAge.touch("Cameron");
}

```

After the last “touch” line above, Cameron’s association will be the most recently written one in the HashTable, followed by Mikey’s, Sally’s, and Lily’s. Timothy’s will not be in the recently-written list, since it is marked as permanent.

Touching a permanent item in the hash table has no effect on the association, since it’s not being tracked in your recently-written list.

Note that only *adding*, *updating* or *touching* a non-permanent association may adjust that associations’s position in the recently-written list. Simply looking up an association in the hash table, for example, **must not** change the association’s position in the recently-written list.

Why would you want to track when each item in the *HashTable* was last written? Well, it is possible that the client of your hash table will eventually insert a large number of associations and the table will reach its capacity. In this situation, the client may wish to discard one or more old associations from the hash table to make room for new ones. By maintaining a log of which items were least recently written, you can easily allow the

client to discard old associations to make room for new ones. Your *HashTable* class template must therefore provide a *discard()* method for this purpose.

Each time it is called, the *discard()* method must discard the least recently written non-permanent association in the hash table (removing it completely), assuming there is at least one non-permanent association in the hash table. Here's how it might be used:

```
unsigned int computeHash(std::string s)
{
    // Don't forget to write this!
}

void argh()
{
    HashTable<string, int> nameToAge(100, 200);

    nameToAge.set("Carey", 43);
    nameToAge.set("David", 97);
    nameToAge.set("Timothy", 43, true);
    nameToAge.set("Ivan", 28);
    nameToAge.set("Sally", 22);
    nameToAge.set("David", 55);

    nameToAge.touch("Carey");

    // let's discard the two least recently written items

    for (int k = 0; k < 2; k++)
    {
        string discardedName;
        int discardedAge;
        if (nameToAge.discard(discardedName, discardedAge))
            cout << "Discarded " << discardedName
                  << " who was " << discardedAge
                  << " years old.\n";
        else
            cout << "There are no items to discard!\n";
    }
}
```

The first loop iteration will discard Ivan's association, since his was the least recently written one in the hash table. The second loop iteration would discard Sally's association, since after Ivan's was discarded, hers was the least recently written one.

After the loop is completed, the hash table holds just Timothy's, David's, and Carey's associations, with Carey's being the most recently written one, and David's being the least recently written one. Timothy's, being permanent, would not be in the recently-written list. David's would be the next association discarded if *discard()* were called again.

The HashTable Class Template: Details

Your *HashTable* implementation **must** have the following public interface; you must **not** change or add to the public interface:

```
template <typename KeyType, typename ValueType>
class HashTable
{
public:
    HashTable(unsigned int numBuckets, unsigned int capacity);
    ~HashTable();
    bool isFull() const;
    bool set(const KeyType& key, const ValueType& value, bool permanent = false);
    bool get(const KeyType& key, ValueType& value) const;
    bool touch(const KeyType& key);
    bool discard(KeyType& key, ValueType& value);

private:
    // We prevent a HashTable from being copied or assigned by declaring the
    // copy constructor and assignment operator private and not implementing them.
    HashTable(const HashTable&);
    HashTable& operator=(const HashTable&);
};
```

Here are the general requirements for your *HashTable* class:

1. You **must** implement your own open hash table in your *HashTable* class template (i.e., define your own *Node* struct/class, maintain an array of pointers to nodes, etc.).
2. You **must not** use any STL containers to implement *HashTable* (e.g., no map, set, unordered_map, unordered_set, vector, list, queue, stack, etc.).
3. Your *HashTable* class template **must** use the public interface documented above. You may add only private members; you must **not** add other public members to *HashTable*. Doing so will result in a score of **zero** for this part of the project.
4. You may assume that the key type of any instantiation of the *HashTable* class template has the comparison operators == and != defined for it (certainly ints and strings do). If strings are keys, they are case-sensitive, so “bill” and “Bill” would be different keys.
5. You **must** define one non-member *computeHash* function for every type of key you wish to store in a hash table. These non-member functions should be implemented in one of your .cpp files.

Requirements for HashTable(unsigned int numBuckets, unsigned int capacity)

The constructor must allocate and initialize an empty hash table with the specified number of buckets. It must also remember the capacity of the hash table. It must run in $O(B)$ time, where B is numBuckets.

Requirements for `~HashTable()`

The destructor must free all memory associated with the hash table's nodes. This method must run in $O(B+N)$ time, where N is the number of items in the hash table and B is the number of buckets in the hash table.

Requirements for `bool isFull() const`

The *isFull()* method must return *true* if the hash table is full to capacity (i.e., the number of associations in the hash table is equal to the capacity), otherwise *false*.

Requirements for `bool set(const KeyType& key, const ValueType& value, bool permanent = false)`

Let B be the number of buckets in the hash table, and C be its capacity. If the indicated key is not already in the table: If the table already contains C associations, the *set()* method returns *false* without changing anything; otherwise it adds a new association $\text{key} \rightarrow \text{value}$, which will be permanent if the third argument is *true*, otherwise non-permanent.

If an association with the indicated key is already in the table, it is updated with the indicated value replacing the value currently stored in that association. In this case, the third argument is ignored; the permanent/non-permanent status of the association may or may not be what is indicated by the third parameter, and is not, for example, changed because of it.

Each non-permanent, newly added or updated association must be made the most recently written item in the recently-written list. Permanent associations must not be tracked in the recently-written list. The method return *true* if an association was added or updated.

Attempting to add or update an association must run in $O(C/B)$ time, as must updating the recently-written list.

Requirements for `bool get(const KeyType& key, ValueType& value) const`

If the indicated key is not in the table, the *get()* method returns *false* and the second parameter is unchanged. Otherwise, the second parameter is set to the value associated with the key in the table, and the method returns *true*. This method does **not** adjust the order of any items within the recently-written list.

If B is the number of buckets in the hash table, and C is its capacity, this method must run in $O(C/B)$ time.

Requirements for `bool touch(const KeyType& key)`

If a non-permanent association in the table has the indicated key, the *touch()* method moves it to the top of the recently-written list, just as if the association

had actually been updated, and returns *true*. Otherwise, it does nothing and returns *false*.

If B is the number of buckets in the hash table, and C is its capacity, this method must run in $O(C/B)$ time.

Requirements for `bool discard(KeyType& key, ValueType& value)`

If there are no non-permanent associations in the table, the *discard()* method does nothing (leaving its parameters unchanged) and returns *false*. Otherwise it must identify the least recently written non-permanent association in the table. It must set the key and value parameters to the key and value of that association, delete that association from the table (ensuring that the recently-written list is still consistent so that future calls to *discard()* will properly discard the next least-recently used association), and return *true*.

If B is the number of buckets in the hash table, and C is its capacity, this method must run in $O(C/B)$ time.

Requirements for your *computeHash()* functions

So what are these *computeHash()* functions that you need to define?

Well, as we learned in class, when using a hash table, you need some way of determining in which bucket to find the information associated with a given key. This is done with a hash function.

The hash function is expected to take a key as its argument and produce an unsigned integer as its result. That result can then be used to compute the bucket number in the hash table.

If you're using several hash tables, you must have a different hash function, each named *computeHash()*, for every *type* of key you will be using in those hash tables. If a application needs two hash tables that use ints as keys, five that use strings as keys, and one that uses *Rectangles* as keys, you'll need three *computeHash()* functions, one for each type. The five hash tables that use strings as keys will all use the same *computeHash()* function, the one for strings.

As an example, for this *HashTable* that maps *Rectangle* objects to *Colors*:

```
HashTable<Rectangle, Color> rectToColor;
```

you'd need to define a hashing function that is capable of computing an unsigned integer value for given a *Rectangle*; here's one possibility:

```
unsigned int computeHash(Rectangle r)
{
```

```

        unsigned int hashCode = r.getX() + r.getY() +
                                r.width() + r.height();
        return hashCode;
    }

```

If you wanted to have a *HashTable* mapping *Student* objects to GPAs (a double):

```

HashTable<Student, double> gpas;

```

you’d need to define a hashing function that is capable of computing an unsigned integer value for each *Student*; here’s one possibility:

```

unsigned int computeHash(Student stud)
{
    return stud.getID(); // use student ID as hash value
}

```

These *computeHash()* non-member functions should be implemented in one of your .cpp files.

When you write your *HashTable* class template, you can then assume that a client of your *HashTable* will supply a suitable *computeHash()* function for the type of key the table will use. You can use it to determine which bucket to put a new association into. Here’s a little hint how this might be used:

```

template<typename KeyType, typename ValueType>
class HashTable
{
    ...
    unsigned int getBucketForKey(const KeyType& key)
    {
        // The computeHash function must be defined
        // for each type of key that we use in some
        // hash table.
        unsigned int computeHash(KeyType); // prototype
        unsigned int result = computeHash(key);
        unsigned int bucketNum =
            determineBucketFromHash(result);
        return bucketNum;
    }
    ...
};

```

By using this approach, you enable your *HashTable* class to work the same way no matter what C++ type is being used as the key. **Important note:** Before calling *computeHash()* in a *HashTable* member function, make sure that the member function has declared a prototype for the *computeHash()* function. While some compilers (e.g., Visual C++ 2013, older versions of g++) are forgiving in some circumstances, Standard C++ has arcane “two-phase lookup” rules for names used in templates that often surprise even experienced C++ programmers.

All hash values for keys **must** be computed using an appropriate *computeHash()* function, which **must** have a prototype in the following form, where *sometype* is the appropriate key type:

```
unsigned int computeHash(sometype key);
```

This will enable us to test your class with our test code, and let us use your hash table with types that it's never seen before. Failure to properly implement this requirement will result in a **zero** on this part of the project.

BinaryConverter Class

The *BinaryConverter* class provides two services:

1. Convert a vector of unsigned short integers into a C++ string of tab and space characters that represent the binary encoding of the input numbers.
2. Convert a C++ string of tab and space characters that represent a binary encoding into a vector of unsigned short numbers that encoding represents.

In this project, you'll be encoding a message like "Meet me at Ackerman at noon" and embedding the encoding into a web page. This will be done in four steps:

1. First you'll convert your message from a C++ string into a series of unsigned short values, e.g. {77, 101, 101, 116, 32, ...}. We'll tell you how to do this later.
2. Then you'll convert these unsigned short values into their binary equivalents of 1s and 0s {0000000001001011, 0000000000110101, 0000000001100101, ...}.
3. Then you'll convert the 1s and 0s to a string containing tabs and spaces (a tab representing a 1, a space representing a 0).
4. Finally, you'll insert these tabs and spaces into a file at the end of each line. These spaces and tabs will covertly encode your secret message within the web page.

The *BinaryConverter* class will be used for steps 2 and 3 above. Similarly, when you want to extract and decode a secret message from a web page, you'll use this class to convert the tabs and spaces from the page back into binary, and from there into their unsigned short form.

You must implement two static public member functions for the *BinaryConverter* class: *encode()* and *decode()*. As shown below, a static member function of a class is not called on a particular object of that class type; it's called through the class name itself. (Unlike the ordinary non-static member functions we've been using all along, a static member function is not passed a *this* pointer.)

Here's an example of how the *encode()* method can be used:

```
void encodeVectorAsBitString()
{
    vector<unsigned short> v;
```

```

v.push_back(1);
v.push_back(5);

// We use the :: operator below because encode() is a static
// member function. You don't create a BinaryConverter object
// to use encode(); instead, you call it using the class name.

string hiddenMessage = BinaryConverter::encode(v);
cout << hiddenMessage; // prints tabs and spaces
}

```

If we were to run the above function, the *hiddenMessage* variable would hold the following string (depicted here with hyphens representing tabs and underscores representing spaces):

-----'------'-'

This represents the following binary string, which ultimately represents {1, 5}:

00000000000000000100000000000000101

Here's an example of how you might use *decode()*:

```

void decodeBitStringAsVector()
{
    vector<unsigned short> v;
    // Using the -/_ representation in this comment, the string
    // below contains -----'------'-
    string msg = "          \t          \t \t";

    if (BinaryConverter::decode(msg, v))
    {
        cout << "The vector has " << v.size() << " numbers:";
        for (int k = 0; k != v.size(); k++)
            cout << ' ' << v[k];
        cout << endl;
    }
    else
        cout << "The string has a bad character or a bad length.";
}

```

If we were to run the above, the output would be

The vector has 2 numbers: 1 5

Most of you haven't learned binary encoding yet (you will in CS 33), so in the provided code, we'll give you two functions, *convertNumberToBitString()* and *convertBitStringToNumber()*, to help you convert an unsigned short value to a binary string of 1s and 0s, and vice versa.

Our *convertNumberToBitString()* function takes an unsigned short as a parameter and returns a C++ string of ‘1’ and ‘0’ characters that represents the input number:

```
string bitString = convertNumberToBitString(19);
cout << bitString;    // writes 0000000000010011
```

The *convertBitStringToNumber()* function takes a C++ string of ‘1’ and ‘0’ characters as a parameter and returns the unsigned short value that the string represents:

```
string bitString = "0000000000010010";
unsigned short value;
if (convertBitStringToNumber(bitString, value))
    cout << value;          // prints 18
else
    cout << "You passed in a bad bit string!\n";
```

The BinaryConverter Class: Details

Here is the required public interface of the *BinaryConverter* class:

```
class BinaryConverter
{
public:
    static std::string encode(const std::vector<unsigned short>& numbers);
    static bool decode(const std::string& bitString,
                      std::vector<unsigned short>& numbers);
};
```

Requirements for `static std::string encode(const std::vector<unsigned short>& numbers)`

This method must accept a vector containing zero or more unsigned short values and return a string with a multiple of sixteen tabs and spaces, with each tab representing a 1 and each space representing a 0. Each successive block of sixteen characters in the string represents the binary encoding of successive numbers in the vector.

Requirements for `static bool decode(const std::string& bitString, std::vector<unsigned short>& numbers)`

If the string does not contain a multiple of sixteen characters, or contains characters other than tabs and spaces, the method must return *false*. The value of the *numbers* parameter in this case is not defined by this spec. (Thus, you’ll presumably just let its value end up as whatever your code happens to have done up to the time you return *false*: its value might be unchanged, or cleared, or partially filled, or whatever.)

If the string contains a multiple of sixteen characters, each a tab or a space, then this method must convert each successive block of sixteen characters in the string to the number of which that block is the binary encoding, where a tab character represents a 1 and a space character represents a 0, placing those numbers into the vector. This method

must discard any data that was in the vector prior to the time it was called. The method must return *true* to indicate a successful conversion.

Compressor Class

The *Compressor* class provides two services:

1. Compress a string, e.g. “Meet me in Ackerman at noon”, using a compression algorithm to produce the compressed result as a series of unsigned short integers. (Other classes will take these unsigned short integers that represent the compressed message, convert them to tabs and spaces using the *BinaryConverter* class, and then insert them into a web page).
2. Decompress a previously-created vector of unsigned shorts to produce the string that was compressed.

Why would you want to compress the data before encoding it secretly into a web page? Because some messages may be extremely long – for instance, you may wish to steganographically encode a whole 1MB secret document within a web page, and since the document is large, you’ll want to compress it first. So for this project, you’re going to build a compression algorithm like one used in the popular ZIP compression format!

ZIP uses a number of different compression algorithms, but one of the key ones is called Lempel-Ziv-Welch (LZW) compression. For this project, you’ll be implementing an LZW variant.

In a nutshell, LZW compression scans through the to-be-compressed data identifying sequences of characters (e.g., “the”) that are repeated over and over within the data, and then replaces later occurrences of these sequences with compact references to the original sequence in order to save space. The algorithm uses a hash table to efficiently discover these repeated sequences and replace them.

For example, at a high level, the input:

the cat eats the bat in the hat

conceptually is replaced by something like:

the cat e2s 1 b2 in 1 h2

where the integers refer to previously seen sequences. The first sequence later repeated is “the”, which the 1 refers to, and the second sequence later repeated is “at”, which the 2 refers to. As you can see, the compressed message is shorter than the first message because repeated sequences like “the” and “at” that occurred earlier in the message are simply replaced by numbers that refer to the original sequence. This is a simplification, but we hope you get the general idea.

So how does our simplified LZW compression algorithm work? Here's the pseudo-code:

To compress a string of characters of length N into a series of unsigned short values:

1. Create a hash table H that maps strings to unsigned shorts. The hash table must have a capacity of C items, where C is the smaller of $N/2 + 512$ (using integer division) and 16384.

The maximum load factor of your hash table must be 0.5. You must determine the number of buckets in the hash table based on the capacity C and this required maximum load factor.

2. Next, use a loop to add 256 associations to your hash table. For the j^{th} association, $0 \leq j \leq 255$, associate a C++ string that contains the single character `static_cast<char>(j)`, with an unsigned short with the same value j . Make sure all of these items are marked as *permanent* when they are inserted with the `set()` method. The effect is as if you did this for 256 characters:

```
...
H.set(" ", 32, true);      // space is mapped to 32
...
H.set("A", 65, true);      // A is mapped to 65
...
H.set("b", 98, true);      // b is mapped to 98
...
```

but of course you'll write a simple loop. After this step, your hash table should hold 256 permanent associations mapping every possible one-character string to a number between 0 and 255 (e.g., "A" \rightarrow 65).³

3. Let `nextFreeID = 256`.
4. Create an empty string named `runSoFar`.
5. Create an empty vector V of unsigned shorts that will contain our compressed result.
6. For each character c in the input string that you're compressing, do the following:
 - a. Create a string named `expandedRun` and set it equal to the `runSoFar` string plus the new character c appended on the end.
 - b. If the `expandedRun` string is already contained in H , then (this longer string was already added to the hash table, so let's keep expanding our run)
 - i. Set `runSoFar` equal to `expandedRun`, increasing the length of the current run by one character.
 - ii. Go on to the next iteration of step 6 above.
 - c. Otherwise our new string `expandedRun` was not found in the hash table, so look up `runSoFar` in hash table H ⁴ and determine what unsigned short it maps to. Call this mapped-to value x .

³ C++ `std::string` objects are allowed to contain zero bytes; one way to create a C++ string of length 1 containing (only) a zero byte is `string s(1, '\0')`; another is `string s(1, static_cast<char>(0))`.

⁴ The proof that it must be in the table is left to the reader.

- d. Append x to the vector V of unsigned shorts.
 - e. Use the hash table's *touch()* method to record the *runSoFar* entry in the hash table as if it were the most recently written association (if it wasn't a permanent association).
 - f. Reset *runSoFar* to the empty string.
 - g. Lookup c in hash table H and determine what unsigned short it maps to. Call this mapped-to value cv .
 - h. Append cv to the vector V of unsigned shorts.
 - i. Now it's time to add our new *expandedRun* sequence to our hash table. If the table is not yet full to capacity, then:
 - i. Associate *expandedRun* in hash table H with the current value of *nextFreeID*.
 - ii. Increment *nextFreeID* by one.
 - j. Otherwise, the table is at capacity, and we need to discard the least-recently-written item in the table and then add our new item to the hash table, reusing the existing ID from the discarded item:
 - i. Use the hash table's *discard()* method to discard the least recently written non-permanent item in the table. Make sure to get the unsigned short value that the just-discarded string was mapped to.
 - ii. Add a new association to the hash table, mapping *expandedRun* to the just-discarded association's unsigned short value.
7. After you have completed step 6 for all the characters in the input string, if the *runSoFar* variable is not empty, then:
- a. Look up *runSoFar* in hash table H and determine what unsigned short it maps to. Call this mapped-to value x .
 - b. Append x to the vector V of unsigned shorts.
8. Finally, append the hash table's capacity C to vector V as the last item. This enables the decompressor to know how big to make its hash table when decompressing the message, since it's critical that its table have the same capacity as that of the compressor.

All right, you're probably sufficiently confused and frightened now, so let's run through a simple example. Imagine that your input string is

AAAAAAAAAB

Since the message length is $N=10$, our capacity C would be 517, and the number of buckets in the table would be 1034.

Now, let's run through steps 6 and 7 for the string above; the hash table starts off with the 256 permanent associations:

- After the first 'A' is processed, *runSoFar* will be "A" and $V = \{\}$
- After the second 'A' is processed, *runSoFar* will be "", $V = \{65, 65\}$, and "AA" → 256 will be added to the hash table.
- After the third A is processed, *runSoFar* will be "A" and $V = \{65, 65\}$

- After the fourth A is processed, *runSoFar* will be “AA” and $V = \{65, 65\}$
- After the fifth A is processed, *runSoFar* will be “”, $V = \{65, 65, 256, 65\}$ and “AAA” \rightarrow 257 will be added to the hash table.
- After the sixth A is processed, *runSoFar* will be “A”, and $V = \{65, 65, 256, 65\}$
- After the seventh A is processed, *runSoFar* will be “AA”, and $V = \{65, 65, 256, 65\}$
- After the eighth A is processed, *runSoFar* will be “AAA”, and $V = \{65, 65, 256, 65\}$
- After the ninth A is processed, *runSoFar* will be “”, $V = \{65, 65, 256, 65, 257, 65\}$, and “AAAA” \rightarrow 258 will be added to the hash table.
- After the B is processed, *runSoFar* will be “B” and $V = \{65, 65, 256, 65, 257, 65\}$
- Since there are no more characters in the input string, we look up the contents of *runSoFar* and append the result, 66, to V: $V = \{65, 65, 256, 65, 257, 65, 66\}$
- Finally, we append the capacity of our hash table, 517, to V: $V = \{65, 65, 256, 65, 257, 65, 66, 517\}$

The final contents of the vector V represent the compressed encoding of the input string. Every integer from 0 to 255 directly encodes one character (e.g., 65 encodes ‘A’). Integers greater than 255 refer to longer strings of characters, such as “AA” or “AAA”, that are discovered and added to the hash table as the algorithm proceeds through the input message. This set of integers exactly corresponds to the original contents of the input string. And as we’ll see below, we can decompress these 8 numbers to reconstruct our original string.

You’ll notice that while our initial string had 10 characters, there are only 8 integers in our vector. Unfortunately, in this short example, since each integer requires 2 bytes (as an unsigned short), the data in our vector actually occupy 16 bytes vs. 10 bytes in our input string. However, if we had compressed a much larger string with lots of repetition (as in a typical document), the overall number of bytes occupied by the vector data would be significantly less than that of the input string.

Now let’s examine an example that adds more items to the hash table than its capacity allows for, and see what is produced. Let’s assume that the hash table has a capacity of 257 (hypothetically, since this the spec’s formula in step 1 of the compression algorithm implies a capacity of at least 512). How would the above compression change?

Again, let’s run through steps 6 and 7 for the string AAAAAAAAAAB; the hash table starts off with the 256 permanent associations:

- After the first ‘A’ is processed, *runSoFar* will be “A” and $V = \{\}$
- After the second ‘A’ is processed, *runSoFar* will be “”, $V = \{65, 65\}$, and “AA” \rightarrow 256 will be added to the hash table, so the table now contains 257 associations.
- After the third A is processed, *runSoFar* will be “A” and $V = \{65, 65\}$
- After the fourth A is processed, *runSoFar* will be “AA” and $V = \{65, 65\}$

- After the fifth A is processed, *runSoFar* will be "", $V = \{65, 65, 256, 65\}$ and we need to add a new association. Because the hash table now contains 257 associations, its capacity, the least recently written (and only) non-permanent association, "AA" \rightarrow 256, will be replaced by "AAA" \rightarrow 256 in the hash table.
- After the sixth A is processed, *runSoFar* will be "A", and $V = \{65, 65, 256, 65\}$
- After the seventh A is processed, *runSoFar* will be "", $V = \{65, 65, 256, 65, 65, 65\}$, and "AAA" \rightarrow 256, the least recently written non-permanent item, will be replaced by "AA" \rightarrow 256 in the hash table.
- After the eighth A is processed, *runSoFar* will be "A", and $V = \{65, 65, 256, 65, 65, 65\}$
- After the ninth A is processed, *runSoFar* will be "AA", and $V = \{65, 65, 256, 65, 65, 65\}$
- After the B is processed, *runSoFar* will be "" and $V = \{65, 65, 256, 65, 65, 65, 256, 66\}$
- Since there are no more characters in the input string, and *runSoFar* is empty, we are done compressing.
- Finally, the total capacity of our hash table, 257, is appended to V: $V = \{65, 65, 256, 65, 65, 65, 256, 66, 257\}$

Notice that our new output sequence had 9 numbers instead of the 8 numbers of the previous example. This is because our hash table had a smaller capacity and was therefore unable to keep track of as many of the original multi-character sequences and replace reoccurrences of these sequences with encoded values. There's often a tradeoff between hash table size and the compression ratio that's achieved. This is a tradeoff that every compression program has to make – use more memory and achieve better compression, or use less memory but achieve worse compression.

OK, so we learned how to compress data, now how about decompressing it? Here's the algorithm:

To decompress a previously-created vector of unsigned short values into the original message string:

1. Create a hash table H that maps unsigned shorts to strings. The hash table must have a capacity equal to the last value in the input vector.

The maximum load factor of your hash table must be 0.5. You must determine the number of buckets in the hash table based on the capacity and this required maximum load factor.

2. Next, similar to what is done in the compression algorithm, add 256 associations to your hash table. Associate each unsigned short j , $0 \leq j \leq 255$, with a C++ string that contains the single character `static_cast<char>(j)`. Make sure all of these items are marked as *permanent* when they are inserted with the `set()` method. The effect is as if you did this for 256 characters:


```

...
H.set(32, " ", true);          // 32 is mapped to space
...
H.set(65, "A", true);          // 65 is mapped to A
...
H.set(98, "b", true);          // 98 is mapped to b
...

```

but of course you'll write a simple loop. After this step, your hash table should hold 256 permanent associations.

3. Let *nextFreeID* = 256.
4. Create an empty string named *runSoFar*.
5. Create an empty string named *output* that will contain our decompressed result.
6. For *us* taking on the value of each of the first N-1 unsigned shorts in the input vector that you're trying to decompress, do the following:
 - a. If *us* is less than or equal to 255 then it represents a normal 1-byte character:
 - i. Look up the (1-character) string associated with *us* in the hash table and append it to the *output* string.
 - ii. If the *runSoFar* string is currently empty, then update its value so it contains a single character, the string associated with *us*, and advance to the next item in the vector by going on to the next iteration of step 6 above.
 - iii. Otherwise, if the *runSoFar* string is non-empty, then create a string named *expandedRun* and set it equal to the *runSoFar* string plus the string associated with *us* appended to it.
 - iv. If the hash table is below capacity, then:
 1. Add a mapping from *nextFreeID* to the value of *expandedRun*.
 2. Increment *nextFreeID*.
 - v. Otherwise the hash table is at capacity, so:
 1. Find and *discard()* the least recently written non-permanent item in the hash table. Make sure to get the unsigned short key of this association.
 2. Add a new association to the hash table, mapping the key of the just discarded item to *expandedRun*.
 - vi. Reset *runSoFar* to the empty string.
 - vii. Go on to the next iteration of step 6.
 - b. Otherwise, *us* represents a multi-character string, not a single character, so:
 - i. If the hash table does not hold an association whose key is *us* then there is an error. Immediately return with a failure indication.
 - ii. Otherwise it does hold such an association, so *touch()* that association to record it as if it were the most recently written one in the hash table (if it wasn't a permanent one).
 - iii. Append the string *S* associated with *us* to the *output* string.
 - iv. Set *runSoFar* to the value of *S*.
 - v. Go on to the next iteration of step 6.

7. Once you have completed step 6 for the first N-1 unsigned shorts in the vector, your *output* string will hold the decompressed message!

So, let's use this algorithm to decode the vector that we created in the most recent example above:

V = {65, 65, 256, 65, 65, 65, 256, 66, 257}

We'll start by creating our new hash table, with a capacity of 257 (using the last value in the vector) and 514 buckets. We continue:

- After the first 65 is processed, *output* will be "A" and *runSoFar* = "A"
- After the second 65 is processed, *output* will be "AA", *runSoFar* = "", and the hash table will hold the new mapping {256 → "AA"}
- After the 256 is processed, *output* will be "AAAA" and *runSoFar* = "AA", with the hash table mapping 256 → "AA" becoming the most recently written item (via *touch()*).
- After the next 65 is processed, *output* will be "AAAAA", *runSoFar* = "", and the hash table will have its {256 → "AA"} mapping changed to {256 → "AAA"}
- After the next 65 is processed, *output* will be "AAAAAA" and *runSoFar* = "A"
- After the next 65 is processed, *output* will be "AAAAAAA", *runSoFar* = "", and the hash table will have its {256 → "AAA"} mapping changed to {256 → "AA"}
- After the next 256 is processed, *output* will be "AAAAAAAAA" and *runSoFar* = "AA", with your hash table mapping 256 → "AA" becoming the most recently written item (via *touch()*).
- After the 66 is processed, *output* will be "AAAAAAAAAAB", *runSoFar* = "", and the hash table will have its {256 → "AA"} mapping changed to {256 → "AAB"}

Notice that the final 257 is not processed after the 66, since it is not part of the compressed message. Your decompression algorithm is now complete, with your *output* string "AAAAAAAAAAB" matching the original input string!

The Compressor Class: Details

Here is the required public interface of the *Compressor* class:

```
class Compressor
{
public:
    static void compress(const std::string &s,
                        std::vector<unsigned short>& numbers);
    static bool decompress(const std::vector<unsigned short>& numbers,
                          std::string& s);
};
```

Here are the requirements for the *Compressor* class:

1. It **must** use *HashTable* objects in its implementation.
2. Except for *vector* and *string*, it **must not** use any STL containers in its implementation (i.e., no *map*, *set*, *unordered_map*, *unordered_set*, *list*, *queue*, *stack*, etc.).

Requirements for `static void compress(const std::string& s, std::vector<unsigned short>& numbers)`

This method must accept a string containing the text you wish to compress. When this function returns, the vector parameter must hold a sequence of unsigned short values encoded as described in the section above. This method must discard any data that was in the vector prior to the time it was called. (It is possible for *s* to be empty, in which case *numbers* will end up containing just a hash table capacity.)

This method must run in $O(L)$ time, where *L* is the length of the string *s* that is to be compressed.

Requirements for `static bool decompress(const std::vector<unsigned short>& numbers, std::string& s)`

This method must accept a vector that is the encoding of a string as described in the section above and decompress the contents of this vector into the string parameter. If there is a decompression error, this method must indicate the failure by returning *false* (in which case the value of the *s* parameter is not defined by this spec – it is allowed to be anything); otherwise the method must return *true*, with the *s* parameter containing the decompressed string. (It is possible for *numbers* to contain just a hash table capacity, in which case *s* will end up empty. If *numbers* is empty, the decompression fails for lack of a hash table capacity.)

This method must run in $O(L)$ time, where *L* is the length of the result string *s*.

Steg Class

The *Steg* class provides two services:

1. Given a string containing the text of a web page (e.g., “<html>This is\nmy page.</html>”) and an input secret message (e.g., “Meet me in Ackerman at noon”), produce a string resulting from compressing and encoding the secret message and embedding it into the web page string.
2. Given a host string with a hidden message, reveal that message by extracting the encoded message from the host string and decoding and decompressing it to produce the secret message.

First, we consider the *hostIn* string to be divided into lines separated by newline characters (`'\n'`). **Caution:**

Substring 1: _____	Substring 4: _____	Substring 7: ____-
Substring 2: -__-	Substring 5: __--	Substring 8: --__
Substring 3: -__-	Substring 6: -__-	Substring 9: --__

If you remember the big picture, you can guess what comes next: We consider each of these N substrings to be appended to the corresponding stripped line:

Line 1: <html>_____	Line 4: _____	Line 7: ____-
Line 2: Q_ -QQ-__-	Line 5: GG__--	Line 8: DDD--__
Line 3: BBB-__-	Line 6: BBB-__-	Line 9: EEE_</html>--__

The *hostOut* string we produce is the concatenation of these N pieces with a newline after each one. Using `|` to represent a newline, this is

```
<html>____|Q_ -QQ-__-|BBB-__-|____|GG__--|BBB-__-|__-|DDD--__|EEE_</html>--__|
```

Notice that if the encoding of *msg* has a length L that is smaller than N , the number of lines in *hostIn*, the rule above requires that the first $L\%N=L$ substrings are of length $(L/N)+1=1$, while the rest (i.e., the last $N-L$) are of length 0, which makes sense. The rule is also reasonable when L is a multiple of N , resulting in all substrings being of length L/N .

Going from hiding to revealing, here is the algorithm that given a *host* string that contains a message hidden by the hiding algorithm above, produces the message:

First, we consider the *host* string to be divided into lines separated by newline characters. The same caution as in the hiding algorithm applies: The last line might not contain a newline, and line separators might be `"\r\n"` sequences instead of just newlines. As an example, if the *host* string is

```
<html>____|Q_ -QQ-__-|BBB-__-|____|GG__--|BBB-__-|__-|DDD--__|EEE_</html>--__|
```

we'd consider it to be divided into 9 lines:

Line 1: <html>_____	Line 4: _____	Line 7: ____-
Line 2: Q_ -QQ-__-	Line 5: GG__--	Line 8: DDD--__
Line 3: BBB-__-	Line 6: BBB-__-	Line 9: EEE_</html>--__

We then gather the trailing tabs and spaces from the lines into one string:

```
_____-__-_-_____-__-_-__-_-__-_-
```

Now we decode that string and decompress the result into a final string (using the services of the *BinaryConverter* and *Compressor* classes, of course). That final string is our result.

The Steg Class: Details

Here is the required public interface of the *Steg* class:

```
class Steg
{
public:
    static bool hide(const std::string& hostIn, const std::string& msg,
                    std::string& hostOut);
    static bool reveal(const std::string& host, std::string& msg);
};
```

Requirements for `static bool hide(const std::string& hostIn, const std::string& msg, std::string& hostOut)`

If *hostIn* is empty, this method returns *false*, leaving *hostOut* unchanged. Otherwise, given a (possibly empty) message to hide, this method must compress *msg* into a sequence of unsigned shorts, encode the unsigned shorts as a string of tabs and spaces that represent the compressed message in binary, produce *hostOut* by appending the tabs and spaces to the ends of each stripped line of the *hostIn* string as described in the section above, and return *true*.

Requirements for `static bool reveal(const std::string& host, std::string& msg)`

This method must extract the tabs and spaces from the end of each line in the *host* string, concatenate them all together, decode the resulting string into a sequence of unsigned shorts, and decompress that sequence into *msg*. If this succeeds, the method returns *true*; if this fails, the method returns *false*. (Failure can occur because the tabs and spaces, if any, at the ends of the lines of *host* were not deliberately placed there by our hiding scheme, but just happened to be there for some other reason, or no intended reason at all, and can't possibly represent any encoded message.)

WebSteg Class

The *WebSteg* class provides two services:

1. Given a message and a string containing the URL of a web page, e.g. “http://cs.ucla.edu”, retrieve that page from the Internet, steganographically embed the message within the text of that page, and return the resulting text.
2. Given a string containing the URL of a web page that contains a steganographically-encoded message, retrieve that page from the Internet, extract the secret message from the text of that page, and return it.

The WebSteg Class: Details

Here is the required public interface of the *WebSteg* class:

```
class WebSteg
{
public:
    static bool hideMessageInPage(const std::string& url,
                                const std::string& msg,
                                std::string& host);
    static bool revealMessageInPage(const std::string& url, std::string& msg);
};
```

Requirements for `static bool hideMessageInPage(const std::string& url, const std::string& msg, std::string& host)`

This method must retrieve a web page from the specified URL, then compress, encode and embed *msg* within the text of that page, producing the resulting text in *host*. If this succeeds, this method must return *true*. If the page can't be retrieved or *msg* can't be embedded within the text of the page, this method must return *false*.

Requirements for `static bool revealMessageInPage(const std::string& url, std::string& msg)`

This method must retrieve a web page from the specified URL, then extract, decode and decompress the message embedded within the text of that page, producing the resulting message in *msg*. If this succeeds, this method must return *true*. If the page can't be retrieved or a message can't be decoded, this method must return *false*.

Requirements and Other Thoughts

Make sure to read this entire section before beginning your project!

1. In Visual C++, make sure to change your project from UNICODE to Multi Byte Character set, by going to Project → Properties → Configuration Properties → General → Character Set
2. In Visual C++, make sure to add *wininet.lib* to the set of input libraries, by going to Project → Properties → Linker → Input → Additional Dependencies ; otherwise, you'll get a linker error!
3. The entire project can be completed in roughly 350 lines of C++ code beyond what we've already written for you, so if your program is getting much larger than this, talk to a TA – you're probably doing something wrong.
4. Before you write a line of code for a class, think through what data structures and algorithms you'll need to solve the problem. How will you use these data structures? Plan before you program!

5. Don't make your program overly complex – use the simplest data structures possible that meet the requirements.
6. You must not modify any of the code in the files we provide you that you will not turn in; since you're not turning them in, we will not see those changes. We will incorporate the required files that you turn in into a project with special test versions of the other files.
7. Make sure to implement and test each class independently of the others that depend on it. Once you get the simplest class coded, get it to compile and test it with a number of different unit tests. Only once you have your first class working should you advance to the next class.
8. We'll provide you with a working version of the *HashTable* class template that uses types from the STL that you are forbidden from using in your implementation. You can use this class template to build and test your other classes even if you can't figure out how to implement your *HashTable* class template.
9. You may use only those STL containers (e.g., vector, list) that are not forbidden by this spec. For *HashTable*, this means you must use none at all. For other classes, this means you must not use *map*, *multimap*, *unordered_map*, *unordered_multimap*, or the nonstandard *hash_map*; use your *HashTable* if you need a map, for example.
10. Try your best to meet our big-O requirements for each method in this spec. If you can't figure out how, then solve the problem in a simpler, less efficient way, and move on. Then come back and improve the efficiency of your implementation later if you have time.

If you don't think you'll be able to finish this project, then take some shortcuts. For example, use the substitute *HashTable* we will provide instead of creating your own *HashTable* class if necessary to save time.

You can still get a good amount of partial credit if you implement most of the project. Why? Because if you fail to complete a class (e.g., *HashTable* or *Compressor*), we will provide a correct version of that class and test it with the rest of your program. If you implemented the rest of the program properly, it should work perfectly with our version of the class you couldn't get working, and we can give you credit for those parts of the project you completed.

But whatever you do, make sure that **ALL CODE THAT YOU TURN IN BUILDS** without errors with both Visual Studio and either clang++ or g++!

What to Turn In

You will turn in **six** files:

BinaryConverter.cpp	Contains your binary converter implementation
Compressor.cpp	Contains your compressor implementation

Steg.cpp	Contains your steganographic transformer for strings implementation
WebSteg.cpp	Contains your steganographic transformer for web pages implementation
HashTable.h	Contains your HashTable class template
report.docx, report.doc, or report.txt	Contains your report

You are to declare and implement the *HashTable* class template in HashTable.h. You must implement the other classes' member functions in the appropriate .cpp files. Your implementations of *computeHash()* for int and for string arguments should be in one of the .cpp files. You may add any #includes or constants that you like to these files, as well as support functions for your implementation. Make sure to properly comment your code.

You must submit a brief (You're welcome!) report that describes:

1. Whether any of your classes have bugs or other problems that we should know about.
2. Whether or not each satisfies our big-O requirements, and if not, what you did instead and what the big-O is for your version.
3. How three of your methods work – use high-level pseudocode to describe these:
 - *HashTable's set()* method
 - *HashTable's touch()* method
 - *HashTable's discard()* method

Grading

- 95% of your grade will be assigned based on the correctness of your solution.
- 5% of your grade will be based on your report.

Good luck!